

LAB 18: QUARKUS TOLERANCE POLICIES

Autor: José Díaz

Github Repo: <https://github.com/joedayz/quarkus-bcp-2025.git>

Abre el proyecto **tolerance-policies-start**.

Instructions

This exercise requires you to add resiliency to the `monitor` application. This application is a microservice that gathers monitoring information from cloud instances, such as system information or CPU utilization. The `monitor` application provides data by invoking other microservices, which are simulated, for the sake of simplicity.

- ▶ 1. Open the application project and review the endpoints.
 - 1.1. Navigate to the `~/D0378/tolerance-policies` directory.

```
[student@workstation ~]$ cd ~/D0378/tolerance-policies
```

- 1.2. Open the project with your editor of choice, such as VSCode or vim.
- 1.3. Inspect the endpoints in the `src/main/java/com/redhat/training/MonitorResource.java` file. These endpoints call other services.
 - `/info`
Invokes `InfoService` to get system information about the cloud instance.
 - `/status`
Invokes `StatusService` to get the status of the cloud instance.
 - `/cpu/stats`
Invokes `CpuStatsService` to get CPU data from the cloud instance.
 - `/cpu/predict`
Invokes `CpuPredictionService` to predict the future CPU load of the cloud instance.



► 2. Install the `smallrye-fault-tolerance` extension and start the application.

2.1. Install the `smallrye-fault-tolerance` extension.

```
[student@workstation tolerance-policies]$ mvn \
quarkus:add-extension -Dextension=smallrye-fault-tolerance
...output omitted...
[INFO] [SUCCESS] ... Extension io.quarkus:quarkus-smallrye-fault-tolerance has
been installed
...output omitted...
```

2.2. Start the application in development mode.

```
[student@workstation tolerance-policies]$ mvn quarkus:dev
...output omitted...
[io.quarkus] (...) started in 1.312s. Listening on: http://localhost:8080
...output omitted...
```

► 3. Use the `Retry` policy to make the application resilient to failures in `InfoService`.

3.1. Open a new terminal window and make a request to the `/info` endpoint. The request fails.

```
[student@workstation ~]$ curl localhost:8080/info; echo
{"details":"Error id ...output omitted... }
```

3.2. Inspect the `src/main/java/com/redhat/training/sysinfo/InfoService.java` file. Only one out of five invocations to the `getInfo` method succeed.

3.3. Add the `@Retry` annotation to the `getInfo` method. Set `maxRetries` to 5.

```
@Retry( maxRetries = 5 )
public Info getInfo() {
    ...implementation omitted...
}
```

3.4. Rerun the request and verify that it works.

```
[student@workstation ~]$ curl localhost:8080/info; echo
{"NAME":"Linux","ARCH":"amd64","VERSION":"4.18.0-372.32.1.el8_6.x86_64"}
```

3.5. Inspect the application logs and verify that Quarkus retried the request several times.

```
ERROR [com.red.tra.ser.InfoService] (...) Request #1 has failed
ERROR [com.red.tra.ser.InfoService] (...) Request #2 has failed
ERROR [com.red.tra.ser.InfoService] (...) Request #3 has failed
ERROR [com.red.tra.ser.InfoService] (...) Request #4 has failed
INFO [com.red.tra.ser.InfoService] (...) Request #5 has succeeded
```

► 4. Use the `Timeout` policy to make the application resilient to delays in `StatusService`.

- 4.1. Make a request to the `/status` endpoint. The request takes about five seconds to complete.

```
[student@workstation ~]$ curl localhost:8080/status; echo
Running
```

- 4.2. Inspect the application logs and verify that the request is taking about five seconds to complete.

```
WARN [com.red.tra.sta.StatusService] (...) Request #1 is taking too long...
INFO [com.red.tra.sta.StatusService] (...) Request #1 completed in 5001
milliseconds
```

- 4.3. Inspect the `src/main/java/com/redhat/training/status/StatusService.java` file. Note two aspects:

- The `getStatus` method experiences delays in four out of five invocations.
- The `getStatus` method retries invocations that fail due to a timeout.

- 4.4. Add the `@Timeout` annotation to the `getStatus` method. Throw a timeout error after 200 milliseconds.

```
@Timeout( 200 )
@Retry(maxRetries = 5, retryOn = TimeoutException.class)
public String getStatus() {
    ...implementation omitted...
}
```

- 4.5. Rerun the request and verify that the response is faster.

```
[student@workstation ~]$ curl localhost:8080/status; echo
Running
```

- 4.6. Review the application logs and verify that Quarkus has interrupted the slow invocations and retried them.

```
WARN [com.red.tra.sta.StatusService] (...) Request #1 is taking too long...
WARN [com.red.tra.sta.StatusService] (...) Request #1 has been interrupted after
200 milliseconds
WARN [com.red.tra.sta.StatusService] (...) Request #2 is taking too long...
WARN [com.red.tra.sta.StatusService] (...) Request #2 has been interrupted after
200 milliseconds
WARN [com.red.tra.sta.StatusService] (...) Request #3 is taking too long...
WARN [com.red.tra.sta.StatusService] (...) Request #3 has been interrupted after
200 milliseconds
WARN [com.red.tra.sta.StatusService] (...) Request #4 is taking too long...
WARN [com.red.tra.sta.StatusService] (...) Request #4 has been interrupted after
200 milliseconds
INFO [com.red.tra.sta.StatusService] (...) Request #5 completed in 0 milliseconds
```

- ! — ▶ 5. Use the Fallback policy to make the application resilient to missing data in `CpuStatsService`.

- 5.1. Make a request to the `/cpu/stats` endpoint. The response contains CPU usage time series data. The response also contains the mean and the standard deviation, calculated from the time series data.

```
[student@workstation ~]$ curl -s localhost:8080/cpu/stats | jq
{
  "usageTimeSeries": [
    0.987903386804749,
    0.34275471439780536,
    0.020709840667124446,
    0.35121416539390315,
    0.9863511894860464,
    0.31318722701672885,
    0.7856415790758161,
    0.42147674186562323,
    0.284121753040781
  ],
  "mean": 0.4992622886387308,
  "standardDeviation": 0.319795785721884
}
```

- 5.2. Repeat the request until an error occurs.

```
[student@workstation ~]$ curl -s localhost:8080/cpu/stats | jq
{
  "details": "Error id ..., org.jboss.resteasy.spi.UnhandledException:
java.lang.NullPointerException",
  "stack": "...output omitted..."
}
```

- 5.3. Inspect the application logs. The error occurs when the `getCpuStats` method calls `calculateMean`.

```
WARN [com.redhat.training.cpu.CpuStatsService] (...) Cpu usage data in request #3
contains null values
ERROR [io.quarkus.vertx.http.runtime.QuarkusErrorHandler] (...) HTTP Request to /cpu/
stats failed, error id: ...: org.jboss.resteasy.spi.UnhandledException:
java.lang.NullPointerException
...output omitted...
at
com.redhat.training.cpu.CpuStatsService.calculateMean(CpuStatsService.java:55)
at
com.redhat.training.cpu.CpuStatsService.getCpuStats(CpuStatsService.java:21)
```

- 5.4. Inspect the `src/main/java/com/redhat/training/cpu/CpuStatsService.java` file. One out of three invocations of the `getCpuStats` method fail because the data contains null values. Null values result in an error when the service calculates the mean and the standard deviation.
- 5.5. Add the `@Fallback` annotation to the `getCpuStats` method. Set `getCpuStatsWithMissingValues` as the fallback method.

```
@Fallback( fallbackMethod = "getCpuStatsWithMissingValues" )
public CpuStats getCpuStats() {
    ...implementation omitted...
}
```

5.6. Implement the fallback method. Set the mean and standard deviation values to 0.0.

```
public CpuStats getCpuStatsWithMissingValues() {
    return new CpuStats( series, 0.0, 0.0 );
}
```

5.7. Repeat the request to the /cpu/stats endpoint until you receive a response with null values. The request uses the fallback method and sets the aggregate properties to 0.0.

```
[student@workstation {gls_lab_script}]$ curl -s localhost:8080/cpu/stats | jq
{
  "usageTimeSeries": [
    0.0965490271728523,
    null,
    0.22910115311828105,
    null,
    0.5527746943344609,
    null,
    0.006881053771782275,
    0.22669714994239298,
    0.3583567119779293
  ],
  "mean": 0.0,
  "standardDeviation": 0.0
}
```

- 6. Use the Circuit Breaker pattern to stop sending traffic to the CpuPredictionService when this service becomes unstable.

6.1. Make a request to the /cpu/predict endpoint. The response is the predicted CPU load.

```
[student@workstation ~]$ curl localhost:8080/cpu/predict; echo
0.9822281195867076
```

6.2. Run the predict_many.sh script. This script invokes the /cpu/predict endpoint every second. The requests start failing.

```
[student@workstation ~]$ ~/D0378/tolerance-policies/predict_many.sh
0.4997873140920043
{"details":"Error id 63a4f993-db90-49fe-8c24-24f33...}
{"details":"Error id 63a4f993-db90-49fe-8c24-24f33...}
{"details":"Error id 63a4f993-db90-49fe-8c24-24f33...}
{"details":"Error id 63a4f993-db90-49fe-8c24-24f33...}
```



- 6.3. Press **Ctrl+C** to stop the script.
- 6.4. Inspect the `src/main/java/com/redhat/training/cpu/CpuPredictionService.java` file. The service can only handle one request every two seconds. Otherwise, the service throws an error.
- 6.5. Add the `@CircuitBreaker` annotation to the `predictSystemLoad` method. Set the `requestVolumeThreshold` property to 6, so that the mechanism opens the circuit if three out of six requests fail. Set the `delay` property to 3000, so that the circuit remains open for three seconds.

```
@CircuitBreaker( requestVolumeThreshold = 6, delay = 3000 )
public Double predictSystemLoad() {
    ...implementation omitted...
}
```

- 6.6. Rerun the `predict_many.sh` script. After six requests, most of them failing, the circuit breaker opens the circuit. At this point, the application returns the response `Prediction service is not available at the moment`. The circuit remains open for three seconds, and then the prediction service returns a valid response again.

```
[student@workstation ~]$ ~/D0378/tolerance-policies/predict_many.sh
0.9050798759755502
{"details":"Error id 63a4f993-db90-49fe-8c24-24f33...}
{"details":"Error id 63a4f993-db90-49fe-8c24-24f33...}
{"details":"Error id 63a4f993-db90-49fe-8c24-24f33...}
{"details":"Error id 63a4f993-db90-49fe-8c24-24f33...}
{"details":"Error id 63a4f993-db90-49fe-8c24-24f33...}
Prediction service is not available at the moment
Prediction service is not available at the moment
0.008288100728291115
{"details":"Error id 63a4f993-db90-49fe-8c24-24f33...}
{"details":"Error id 63a4f993-db90-49fe-8c24-24f33...}
{"details":"Error id 63a4f993-db90-49fe-8c24-24f33...}
{"details":"Error id 63a4f993-db90-49fe-8c24-24f33...}
{"details":"Error id 63a4f993-db90-49fe-8c24-24f33...}
Prediction service is not available at the moment
...output omitted...
```

- 6.7. Press **Ctrl+C** to stop the script.

- ▶ 7. Return to the terminal where the application is running in development mode and press **q** to stop the application.

Finish

On the workstation machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish tolerance-policies
```

This concludes the section.

enjoy!

Jose